

A Symmetry No-Go / Go Theorem for Dark-State Interference

in Thermally Broadened Solid-State Emitters

Why room-temperature EIT in the NV center is forbidden, and where it is allowed

Risei Abe

July 7, 2026

Abstract

Electromagnetically induced transparency (EIT) is normally understood through a three-level Λ system in which two long-lived lower states couple to a *single* excited state. In solid-state emitters the optical transition instead addresses a manifold of excited states, and whether a dark state exists is set by the geometry of the probe and control coupling vectors, not by optical linewidths alone. We formulate a basis-invariant criterion built on the coherent interference kernel $K(\omega, T) = \mathbf{p}^\dagger G_e(\omega, T) \mathbf{c}$, where G_e is the non-Hermitian Green's function of the excited manifold. Our central result is that in the thermally merged-manifold limit the kernel collapses to a *temperature-independent* geometric overlap $\mathbf{p}^\dagger \mathbf{c} = \langle g_1 | \hat{\mathbf{d}}_p^\dagger \hat{\mathbf{d}}_c | g_2 \rangle$, whose value is fixed *exactly* by point-group symmetry through a Wigner-Eckart argument. For the NV^- center the two lower states are distinguished only by electron spin within a single orbital, so $\mathbf{p}^\dagger \mathbf{c} = 0$ identically: conventional room-temperature zero-phonon-line spin- Λ EIT is symmetry-forbidden, not merely dephasing-suppressed. Crucially, this is the *opposite* of the naive expectation that phonon merging should build a collective Λ : merging *restores* the bare dipole selection rule that the low-temperature fine structure had broken. The same theorem is constructive: EIT survives merging whenever the two lower states are distinguished by an *orbital* quantum number the dipole can act on with spin overlap intact. We verify this on the inversion-symmetric group-IV centers, where an ab-initio-parametrized SiV^- orbital- Λ ($\lambda_g=61$, $\lambda_e=215$ GHz) reaches $\eta_{\text{geom}} \approx 1$ under cross-polarized drive—the sharp counterpart to the NV^- null—whereas a group-IV *spin*- Λ vanishes for the very same reason NV^- does. The result is a single, symmetry-level design rule for room-temperature coherent optical control in solid-state emitters.

1 Introduction

The three-level Λ picture of EIT rests on a special feature: the probe and control fields drive *the same* excited state, so a coherent superposition of the two lower states can be made dark. A solid-state emitter does not offer a single excited state. The optical transition couples to a manifold $\{|e_j\rangle\}$, and the probe and control address this manifold with different amplitudes *and phases*. The existence of a dark state is then a question about vectors, not levels.

A natural expectation is that raising the temperature should *help*. Once phonon broadening $\Gamma_{\text{opt}}(T)$ exceeds the excited-state fine structure Δ_{ES} , the manifold merges into one featureless optical resonance, and one might hope this recovers an effective single-excited-state Λ . We show that for the NV^- center the opposite is true, and that the reason is a symmetry statement rather than a numerical accident.

It is worth confronting the simplest objection immediately, because it is the objection a referee will raise. One might argue that $\mathbf{p}^\dagger \mathbf{c} = \langle g_1 | \hat{\mathbf{d}}_p^\dagger \hat{\mathbf{d}}_c | g_2 \rangle = 0$ is nothing but the elementary fact that an electric dipole cannot connect two orthogonal *spin* states, and that summing over a complete manifold by closure merely reinstates that bare selection rule. This is correct—and it is precisely the point. At low temperature the NV^- 3E fine structure (spin–orbit and spin–spin interactions) *mixes* spin and orbital character, so the resolved excited states are *not* spin eigenstates; this mixing is exactly what breaks the bare selection rule and enables the well-known cryogenic NV-EIT. Thermal merging removes the energetic weighting that let the system exploit that mixing, and the manifold sum collapses back onto the symmetric closure limit in which the bare rule is restored. The non-trivial content of this paper is therefore not that a spin-only Λ is forbidden in the abstract, but that *broadening converts an allowed low-temperature interference into an exactly forbidden one*, and that the boundary between the two is a computable, sample-tunable curve.

We make three contributions. (i) We recast multi-excited-state EIT through a basis-invariant kernel $K = \mathbf{p}^\dagger G_e \mathbf{c}$ and separate the exact-dark-state condition ($\mathbf{p} \parallel \mathbf{c}$) from finite EIT ($K \neq 0$). (ii) We prove that in the merged limit the surviving overlap $\mathbf{p}^\dagger \mathbf{c}$ is fixed *exactly* by point-group symmetry, giving a no-go/go rule that depends only on which quantum number distinguishes the two lower states. (iii) We apply it to obtain an exact NV^- no-go and an explicit group-IV go, turning a negative NV result into a general, predictive design principle.

2 Basis-invariant EIT criterion

Two long-lived lower states $|g_1\rangle, |g_2\rangle$ couple optically to a manifold $\{|e_j\rangle\}_{j=1}^{N_e}$. In the rotating frame,

$$H_{\text{int}} = -\frac{\hbar}{2} \left(\Omega_p \sum_j p_j |e_j\rangle \langle g_1| + \Omega_c \sum_j c_j |e_j\rangle \langle g_2| + \text{h.c.} \right), \quad (1)$$

with complex coupling vectors $\mathbf{p} = (p_1, \dots, p_{N_e})^T$, $\mathbf{c} = (c_1, \dots, c_{N_e})^T$. The phases are essential: EIT is interference, so transition *strengths* alone are insufficient. The excited manifold responds through the non-Hermitian Green's function

$$G_e(\omega, T) = [\omega \mathbb{I} - H_e(T) + \frac{i}{2} \Gamma_e(T)]^{-1}, \quad (2)$$

where $H_e(T)$ carries the fine structure (spin–orbit, spin–spin, strain, electric- and magnetic-field terms, thermal renormalization) and $\Gamma_e(T)$ the optical decay and phonon-induced broadening. The object that controls transparency is the coherent kernel

$$\boxed{K(\omega, T) = \mathbf{p}^\dagger G_e(\omega, T) \mathbf{c}.} \quad (3)$$

K is invariant under any unitary change of basis in the excited manifold, so it is not an artifact of how one labels the excited states.

In the weak-probe limit the susceptibility is, after a Sherman–Morrison reduction,

$$\chi(\omega, \delta) \propto \frac{i}{2} \left[\mathbf{p}^\dagger D^{-1} \mathbf{p} - \frac{|\Omega_c|^2}{4D_g} \frac{(\mathbf{p}^\dagger D^{-1} \mathbf{c})(\mathbf{c}^\dagger D^{-1} \mathbf{p})}{1 + \frac{|\Omega_c|^2}{4D_g} \mathbf{c}^\dagger D^{-1} \mathbf{c}} \right], \quad (4)$$

with $D_{jj} = \Gamma_j/2 + i(\Delta_j - \omega)$ and $D_g = \gamma_{12}/2 + i\delta$. The control-induced transparency term is proportional to $K_D = \mathbf{p}^\dagger D^{-1} \mathbf{c}$, the diagonal-response form of Eq. (3). Hence a necessary condition

for control-induced transparency through this manifold is $K \neq 0$; if $K = 0$, no value of Ω_c can generate the interference.

This separates cleanly into three regimes:

$$\mathbf{p} \parallel \mathbf{c} \Rightarrow \text{exact dark state}, \quad (5)$$

$$K \neq 0, \mathbf{p} \nparallel \mathbf{c} \Rightarrow \text{finite (imperfect) EIT}, \quad (6)$$

$$K = 0 \Rightarrow \text{no-go for EIT through this manifold}. \quad (7)$$

The exact dark state ($\mathbf{p} \parallel \mathbf{c}$, i.e. $\text{rank}[\mathbf{p} \ \mathbf{c}] = 1$) is strictly stronger than finite EIT and is not required for a usable transparency window.

3 The merged-manifold symmetry theorem

3.1 Collapse to a temperature-independent overlap

Let $\Delta_{\text{ES}} = \max_{j,k} |E_j - E_k|$. When $\Gamma_{\text{opt}}(T) \gtrsim \Delta_{\text{ES}}$ the response no longer resolves individual excited states and $G_e(\omega, T) \simeq g(\omega, T) \mathbb{I}$, so

$$K(\omega, T) \simeq g(\omega, T) \mathbf{p}^\dagger \mathbf{c}. \quad (8)$$

The entire temperature and lineshape dependence factors into the scalar $g(\omega, T)$. What remains is the *geometry*, and the geometry is temperature independent:

$$\mathbf{p}^\dagger \mathbf{c} = \sum_j \langle g_1 | \hat{\mathbf{d}}_p^\dagger | e_j \rangle \langle e_j | \hat{\mathbf{d}}_c | g_2 \rangle = \langle g_1 | \hat{\mathbf{d}}_p^\dagger \Pi_E \hat{\mathbf{d}}_c | g_2 \rangle, \quad (9)$$

where $\Pi_E = \sum_j |e_j\rangle\langle e_j|$ projects onto the merged manifold. This is the key reframing: *the no-go is not a property of the merged lineshape; it is an exact property of the coupling vectors that holds configuration-by-configuration*. Temperature's only role is to strip away the energy denominators $1/(\omega - E_j + i\Gamma_j/2)$ that, at low T , weight the sum unequally and let a resolved subset evade the geometry.

3.2 Symmetry determines the overlap

Because $\hat{\mathbf{d}}$ is an electric dipole—odd under inversion, and *spin-scalar*—the operator $\hat{\mathbf{d}}_p^\dagger \Pi_E \hat{\mathbf{d}}_c$ reduced onto the lower manifold is a spatial tensor of rank ≤ 2 acting trivially on spin. Decompose the lower states into irreducible representations (irreps) of the defect point group G . By the Wigner–Eckart theorem, $\mathbf{p}^\dagger \mathbf{c} \neq 0$ is possible *iff* $|g_1\rangle$ and $|g_2\rangle$ carry the same spin *and* their orbital parts are connected by some irreducible component contained in $\hat{\mathbf{d}}_p^\dagger \hat{\mathbf{d}}_c$. Equivalently:

Merged-manifold no-go/go theorem. In the thermally merged limit, control-induced EIT survives *iff* the two lower states are distinguished by a *spatial* (orbital / polarization-accessible) quantum number. If they differ only in electron spin, $\mathbf{p}^\dagger \mathbf{c} = 0$ exactly and the manifold cannot mediate EIT at any control power.

The polarization dependence is explicit: the tensor components available in $\hat{\mathbf{d}}_p^\dagger \hat{\mathbf{d}}_c$ are selected by $\epsilon_p^* \otimes \epsilon_c$, so probe/control polarization is a physical knob that turns a candidate overlap on or off.

Table 1: Bounds on mechanisms that could break the NV^- merged-manifold no-go. Each is an upper bound on η_{geom} (or on the residual $|\mathbf{p}^\dagger \mathbf{c}|$ envelope). Insert final numerical values from the configuration sweeps.

Mechanism	Bound	Character
Kinematic residual (worst-case envelope, ~ 1 THz)	$\lesssim 1.7 \times 10^{-4}$	lineshape
Hyperfine (^{13}C)	$\lesssim 7.6 \times 10^{-8}$	nuclear-spin
Branch-asymmetric linewidths	$\sim 7 \times 10^{-9}$ (Γ -indep. floor)	kinetic
Spin-orbit-admixed spin-dependent dipole	$\eta_{\text{geom}} \approx 2.5 \times 10^{-10}$	geometric
m_s -diagonal dipole corrections	0 (exact)	symmetry-protected

4 NV^- : an exact no-go

NV^- has C_{3v} symmetry. The optical ZPL connects the ground 3A_2 manifold (orbital *singlet*, a_1 -type; spin triplet $m_s = 0, \pm 1$) to the excited 3E manifold (orbital doublet e). A spin- Λ uses two lower states that share the *same* orbital and differ only in spin. The reduced operator $\hat{\mathbf{d}}_p^\dagger \hat{\mathbf{d}}_c$ on a one-dimensional orbital can only be the scalar (A_1) component times the spin identity, so

$$\mathbf{p}^\dagger \mathbf{c} = (\boldsymbol{\epsilon}_p^* \cdot \boldsymbol{\epsilon}_c) \langle g_1 | g_2 \rangle = 0, \quad \eta_{\text{geom}}^{\text{NV}} = \frac{|\mathbf{p}^\dagger \mathbf{c}|^2}{\|\mathbf{p}\|^2 \|\mathbf{c}\|^2} = 0, \quad (10)$$

identically. There is no orbital tensor component available to rescue the overlap, because the ground orbital is non-degenerate. This is the microscopic dipole sum rule, now seen as a one-line consequence of 3A_2 being an orbital singlet.

Robustness / closed escape routes. The result is leading order in nothing—it is exact within the spin-scalar-dipole, complete-manifold approximation—but the manuscript should quantify every correction that could reintroduce a spatial handle. Table 1 collects the mechanism bounds [values from the microscopic 3E solver; confirm against final runs]. Each candidate either respects the sum rule exactly (spin-diagonal dipole corrections) or contributes a provably tiny geometric overlap.

The physical statement is therefore stronger than “large dephasing kills contrast.” The interference *amplitude* itself vanishes by symmetry when the complete 3E manifold is sampled.

5 Group-IV centers: an explicit go

The theorem is constructive, and the cleanest counterpart is the inversion-symmetric group-IV family (SiV^- , GeV^- , SnV^-), point group D_{3d} , where *both* the ground (2E_g) and excited (2E_u) states are orbital doublets carrying spin 1/2. It is essential to identify the correct Λ . In the merged limit the overlap factorizes,

$$\mathbf{p}^\dagger \mathbf{c} = \langle g_1 | \hat{\mathbf{d}}_p^\dagger \hat{\mathbf{d}}_c | g_2 \rangle = \underbrace{\langle s_1 | s_2 \rangle}_{\text{spin}} \times \underbrace{\langle o_1 | \hat{\mathbf{d}}_p^\dagger \hat{\mathbf{d}}_c | o_2 \rangle}_{\text{orbital}}, \quad (11)$$

because the electric dipole is spin-scalar. A nonzero result needs *both* factors. This immediately rules out the naive choices: a group-IV *spin*- Λ built on the lowest (Kramers) doublet pairs states of *opposite* spin, so $\langle s_1 | s_2 \rangle = 0$ and $\eta_{\text{geom}} = 0$ exactly—just as in NV^- , and despite the orbital




Figure 1: **SiV⁻ satisfies the go condition maximally.** (a) Merged-limit overlap η_{geom} versus probe/control linear-polarization angles for the SiV⁻ orbital- Λ ($\lambda_g=61$, $\lambda_e=215$ GHz): zero for parallel drive, unity for cross-polarization. (b) η_{geom} (optimized cross-pol) versus the ground orbital-branch splitting as strain increases from the ab initio value; pinned near unity, while the NV⁻ spin- Λ stays at zero.

degeneracy. The channel that survives is an *orbital- Λ* : two *same-spin* states in different orbital branches of 2E_g (split by the ground spin-orbit constant λ_g). Then $\langle s_1 | s_2 \rangle = 1$, and the orbital factor is carried by the E_g and A_{2g} components of $\hat{\mathbf{d}}_p^* \hat{\mathbf{d}}_c \subset E_u \otimes E_u = A_{1g} \oplus A_{2g} \oplus E_g$, which connect the two orbital states. The A_{1g} (scalar) piece would only give $\langle o_1 | o_2 \rangle$; it is the non-scalar tensor pieces, selected by $\epsilon_p^* \otimes \epsilon_c$, that activate the channel—maximized by cross-polarized probe and control and vanishing for parallel polarization.

We make this quantitative with ab initio SiV⁻ parameters (Thiering & Gali, PRX **8**, 021063 (2018)): ground/excited spin-orbit splittings $\lambda_g = 61$ GHz, $\lambda_e = 215$ GHz, Huang-Rhys $S = 0.27$. Evaluating Eq. (11) over the complete merged 2E_u manifold (closure verified to $< 10^{-15}$) gives, for the orbital- Λ under cross-polarized in-plane drive,

$$\eta_{\text{geom}}^{\text{SiV}} \simeq 1, \quad (12)$$

i.e. the coupling vectors become collinear ($\mathbf{p} \parallel \mathbf{c}$, the exact-dark-state limit) within the electronic model, while parallel polarization gives the machine-floor null. The result is independent of the one parameter the ab initio dataset does not fix (the axial/in-plane dipole ratio $d_{\parallel}/d_{\perp}$, which in-plane drive never samples) and is stable as strain widens the ground splitting from 61 GHz to $\gtrsim 800$ GHz (Fig. 1). Vibronic and higher excited structure reduce η_{geom} below unity but leave it $\mathcal{O}(1)$; the essential point is that it is bounded *away from zero by symmetry*, in sharp contrast to the exact NV⁻ null. Moreover SiV⁻ emits $\sim 76\%$ into the ZPL ($e^{-S} \approx 0.76$), versus $\sim 3\%$ for NV⁻, so the coherent manifold the theorem addresses carries most of the oscillator strength.

This sharpens the design rule to its precise form:

Design rule. The merged-manifold no-go/go is decided by whether the two lower states are distinguished by a quantum number the dipole can act on *while retaining spin overlap*. Spin-only distinction (NV⁻, orbital-singlet ground; or a group-IV Kramers pair) $\Rightarrow$ *no-go*. Same-spin orbital distinction (group-IV orbital- Λ) $\Rightarrow$ *go*, activated by cross-polarization.

Geometry versus practicality. Passing the geometric test is necessary, not sufficient. The lower states of the orbital- Λ are the two ground orbital branches, split by only $\lambda_g \approx 61$ GHz; their coherence $\gamma_{12} = \gamma_{\text{orb}}(T)$ is single-phonon-limited and degrades above a few kelvin (the known origin of SiV^- thermal decoherence). Whether the power condition $|\Omega_c|^2 \eta_{\text{geom}} \gtrsim \Gamma_{\text{opt}} \gamma_{\text{orb}}(T)$ holds at elevated temperature is a coherence-budget question *separate* from the geometric no-go, and strain—which widens and stiffens the ground splitting against phonons—is the natural route to raise it. The dichotomy is thus: NV^- is excluded at the level of the interference amplitude ($\eta_{\text{geom}} = 0$, exact); SiV^- is permitted geometrically ($\eta_{\text{geom}} \approx 1$, ab initio) and limited only by orbital coherence.

Empirical anchor. That the go side is not vacuous is established experimentally by rare-earth ions (e.g. $\text{Pr}^{3+}:\text{Y}_2\text{SiO}_5$, $\text{Eu}^{3+}:\text{Y}_2\text{SiO}_5$), where robust solid-state EIT and slow light are routine: there the lower states are crystal-field/hyperfine sublevels connected through $4f-4f$ and orbital admixture, so $\mathbf{p}^\dagger \mathbf{c} \neq 0$. The triad *RE (empirical go) / NV (exact no-go) / group-IV (predicted go)* spans the theorem.

6 Escaping the NV no-go: two mechanism classes

For NV^- specifically, the merged-manifold cancellation is escaped only by breaking one of the two hypotheses of the theorem. Every route reduces to one of two classes.

(A) *Prevent equal sampling of the canceling manifold.* Restore an unequal energy weighting so $G_e \neq g\mathbb{I}$ and a non-canceling subset S survives, $\mathbf{p}_S^\dagger \mathbf{c}_S \neq 0$. This is realized by (i) low temperature, $\Gamma_{\text{opt}}(T) \ll \Delta_{\text{ES}}$; (ii) strain splitting $\Delta_{\text{strain}} \gtrsim \Gamma_{\text{opt}}(T)$; (iii) cavity selection of a subset of transitions; (iv) polarization projection $\mathbf{p} \rightarrow P_{\text{pol}} \mathbf{p}$. In all cases the practical EIT condition becomes

$$|\Omega_c|^2 \eta_S \gtrsim \Gamma_{\text{opt}} \gamma_{12}, \quad \eta_S = \frac{|\mathbf{p}_S^\dagger \mathbf{c}_S|^2}{\|\mathbf{p}_S\|^2 \|\mathbf{c}_S\|^2}, \quad (13)$$

i.e. the usual three-level power condition dressed by a geometric efficiency η_S .

(B) *Change the operative manifold.* Use lower states or channels distinguished by a spatial quantum number: singlet-mediated or infrared transitions, Raman/two-photon schemes, or microwave-dressed optical transitions, for which a different kernel $\mathbf{p}_F^\dagger G_F \mathbf{c}_F \neq 0$ applies even though the ZPL spin- Λ is no-go.

Note that both classes are unified by the theorem: they either reintroduce an orbital weighting the dipole can exploit (A) or select lower states with an orbital distinction (B).

7 Falsifiable prediction: the $C(T)$ collapse

The theorem’s decisive, testable consequence is a *non-monotonic* contrast $C(T)$ that rises through the resolved regime, peaks at a strain-dependent T^* , and falls to zero as the 3E manifold merges:

$$C(T) \xrightarrow{\Gamma_{\text{opt}}(T) \gg \Delta_{\text{ES}}} 0, \quad T^*(\text{strain}) \lesssim 30 \text{ K}. \quad (14)$$

This simultaneously (i) reproduces the canonical low-temperature NV-EIT benchmark (Acosta *et al.*: $\sim 6\%$ contrast, sub-MHz windows on the 637 nm ZPL, measured at 6–30 K), and (ii) explains the absence of any room-temperature ZPL-EIT report as a symmetry prohibition rather than a

sensitivity limit. The boundary between the no-go and go regions of the strain–temperature plane is

$$\Delta_{\text{sel}}(\text{strain}) \sim \Gamma_{\text{opt}}(T), \quad (15)$$

which is directly sample-tunable. Two cautions belong in the main text. First, the theorem concerns the coherent ZPL manifold; the vibronic phonon sideband (a few-percent Debye–Waller ZPL fraction at 300 K) is a *different* manifold—either an incoherent background or a separate go-channel—and should be treated as such rather than folded into Γ_{opt} . Second, the merged limit $G_e \simeq g\mathbb{I}$ is precisely the orbital motional-averaging regime of the 3E manifold, which grounds the approximation in established color-center physics rather than an ansatz.

8 Emitter classification and outlook

In the merged limit every emitter falls into one of four classes by $\eta_{\text{geom}} = |\mathbf{p}^\dagger \mathbf{c}|^2 / (\|\mathbf{p}\|^2 \|\mathbf{c}\|^2)$: exact dark-state ($\mathbf{p} \parallel \mathbf{c}$), partial-EIT ($\mathbf{p}^\dagger \mathbf{c} \neq 0$, $\mathbf{p} \nparallel \mathbf{c}$), no-go ($\mathbf{p}^\dagger \mathbf{c} = 0$), and tunable (full-manifold no-go but $\mathbf{p}_S^\dagger \mathbf{c}_S \neq 0$ for a selected submanifold). The classification is a group-theoretic decision on the ground manifold: orbital-singlet ground $\Rightarrow$ spin- Λ no-go; orbital-degenerate ground $\Rightarrow$ go. Computing η_{geom} for group-IV centers, SiC divacancies, rare-earth ions, hBN defects, and quantum dots turns the NV result into a predictive map for room-temperature coherent optical control—the form in which the result is broad enough for a flagship venue.